

Tock Embedded OS Training

RustConf 2017

Please make sure you have completed all of the tutorial pre-requisites. If you prefer, you can download a virtual machine image with all the pre-requisites already installed.

`https://github.com/helena-project/tock/tree/master/doc/courses/rustconf`

Tock is a...

1. Secure
2. Embedded
3. Operating System
4. for Low Resource
5. Microcontrollers

Secure



~~Secure~~ Safe

Tock has isolation primitives that allow you to build secure systems.

Embedded

Definition A

Kernel, applications and hardware are tightly integrated.

Definition B

You'll likely be writing the kernel.

Operating System

Tock itself provides services to components in the system:

- ▶ Scheduling
- ▶ Communication
- ▶ Hardware multiplexing

Low Resource

- ▶ 10s μ A average power draw
- ▶ 10s of kB of RAM
- ▶ Moderate clock speeds

Microcontrollers

System-on-a-chip with integrated flash, SRAM, CPU and a bunch of hardware controllers.


Typically:

- ▶ Communication: UART, SPI, I2C, USB, CAN...
- ▶ External I/O: GPIO, external interrupt, ADC, DAC
- ▶ Timers: RTC, countdown timers


Maybe...

- ▶ Radio (Bluetooth, 15.4)
- ▶ Cryptographic accelerators
- ▶ Other specialized hardware...

Two types of components: capsules and processes



Two types of scheduling: cooperative and preemptive



Part 1: Hardware, tools and development environment

Hail

SAM4L

Accelerometer

MCU


RGB LED

USB

BLE

Temperature &
Humidity

Light Sensor



Binaries on-board

Bootloader

Kernel

Processes

Tools

- ▶ `make` (just instrumenting `xargo`)
- ▶ `Rust` (nightly for `asm!`, compiling core, etc)
- ▶ `xargo` to automate compiling base libraries
- ▶ `arm-none-eabi GCC/LD` to link binaries
- ▶ `tockloader` to interact with Hail and the bootloader

Tools: tockloader

Write a binary to a particular address in flash

```
$ tockloader flash --address 0x1000 \  
    target/thumbv7em-none-eabi/release/hail.bin
```

Program a process in Tock Binary Format¹:

```
$ tockloader install myapp.tab
```

Restart the board and connect to the debug console:

```
$ tockloader listen
```

¹TBFs are relocatable process binaries prefixed with headers like the package name. .tab is a tarball of TBFs for different architectures as well as a metadata file for tockloader.

Check your understanding

1. What kinds of binaries exist on a Tock board? Hint: There are three, and only two can be programmed using `tockloader`.
2. Can you point to the chip on the Hail that runs the Tock kernel? How about the processes?
3. What steps would you follow to program a processes onto Hail? What about to replace the kernel?

Hands-on: Set-up development environment

1. Compile and flash the kernel
2. Compile and program `ble-env-sense` service
3. (Optional) Add some other apps from the repo, like `blink` and `sensors`
4. (Optional) Familiarize yourself with `tockloader` commands
 - ▶ `uninstall`
 - ▶ `list`
 - ▶ `erase-apps`

Part 2: The kernel

Trusted Computing Base (unsafe allowed)

- ▶ Hardware Abstraction Layer
- ▶ Board configuration
- ▶ Event & Process scheduler
- ▶ Rust core library
- ▶ Core Tock primitives

kernel/

chips/

Capsules (unsafe not allowed)

- ▶ Virtualization
- ▶ Peripheral drivers
- ▶ Communication protocols (IP, USB, etc)
- ▶ Application logic

capsules/

Constraints

Small isolation units

Breaking a monolithic component into smaller ones should have low/no cost

Avoid memory exhaustion in the kernel

No heap. Everything is allocated statically.

Low communication overhead

Communicating between components as cheap as an internal function call. Ideally inlined.

Event-driven execution model

```
pub fn main<P, C>(platform: &P, chip: &mut C,
                  processes: &mut [Process]) {
    loop {
        chip.service_pending_interrupts();
        for (i, p) in processes.iter_mut().enumerate() {
            sched::do_process(platform, chip, process;
        }

        if !chip.has_pending_interrupts() {
            chip.prepare_for_sleep();
            support::wfi();
        }
    }
}
```

Event-driven execution model

```
fn service_pending_interrupts(&mut self) {  
    while let Some(interrupt) = get_interrupt() {  
        match interrupt {  
            ASTALARM => ast::AST.handle_interrupt(),  
            USART0 => usart::USART0.handle_interrupt(),  
            USART1 => usart::USART1.handle_interrupt(),  
            USART2 => usart::USART2.handle_interrupt(),  
            ...  
        }  
    }  
}
```


Event-driven execution model

```
impl Ast {
    pub fn handle_interrupt(&self) {
        self.clear_alarm();
        self.callback.get().map(|cb| { cb.fired(); });
    }
}

impl time::Client for MuxAlarm {
    fn fired(&self) {
        for cur in self.virtual_alarms.iter() {
            if cur.should_fire() {
                cur.armed.set(false);
                self.enabled.set(self.enabled.get() - 1);
                cur.fired();
            }
        }
    }
}
```

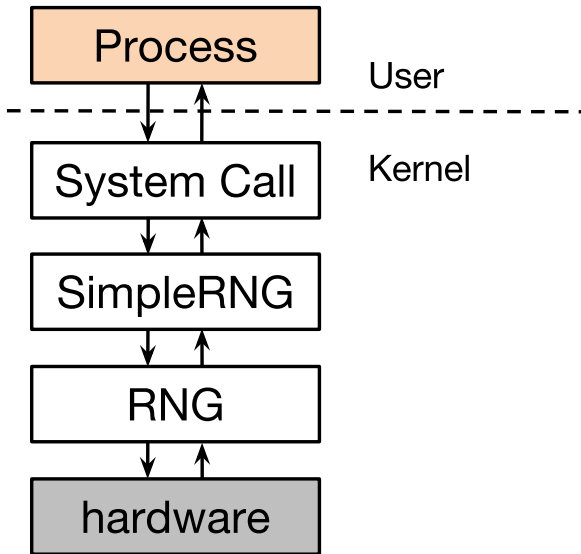


Figure: Capsules reference each other directly, assisting inlining

The mutable aliases problem

```
enum NumOrPointer {  
    Num(u32),  
    Pointer(&mut u32)  
}  
  
// n.b. will not compile  
let external : &mut NumOrPointer;  
match external {  
    Pointer(internal) => {  
        // This would violate safety and  
        // write to memory at 0xdeadbeef  
        *external = Num(0xdeadbeef);  
        *internal = 12345; // Kaboom  
    },  
    ...  
}
```

Interior mutability to the rescue

Type	Copy-only	Mutual exclusion	Opt.	Mem Opt.
Cell	✓	✗	✓	✓
VolatileCell	✓	✗	✗	✓
TakeCell	✗	✓	✗	✓
MapCell	✗	✓	✓	✗

```

pub struct Fxos8700cq<`a> {
    i2c: &`a I2CDevice,
    state: Cell<State>,
    buffer: TakeCell<`static, [u8]>,
    callback:
        Cell<Option<&`a hil::ninedof::NineDofClient>>,
}

impl<`a> I2CClient for Fxos8700cq<`a> {
    fn cmd_complete(&self, buf: &`static mut [u8]) { ... }
}

impl<`a> hil::ninedof::NineDof for Fxos8700cq<`a> {
    fn read_accelerometer(&self) -> ReturnCode { ... }
}

pub trait NineDofClient {
    fn callback(&self, x: usize, y: usize, z: usize);
}

```

Check your understanding

1. What is a `VolatileCell`? Can you find some uses of `VolatileCell`, and do you understand why they are needed?
Hint: look inside `chips/sam4l/src`.
2. What is a `TakeCell`? When is a `TakeCell` preferable to a standard `Cell`?

Hands-on: Write and add a capsule to the kernel

1. Read the Hail boot sequence in `boards/hail/src/main.rs`
2. Write a new capsule that prints “Hello World” to the debug console.
3. Extend your capsule to print “Hello World” every second
4. Extend your capsule to read and report the accelerometer
5. Extra Credit: Write a 9dof virtualization capsule.

Part 3: User space

System calls

Call	Target	Description
command	Capsule	Invoke an operation on a capsule
allow	Capsule	Share memory with a capsule
subscribe	Capsule	Register an upcall
memop	Core	Modify memory break
yield	Core	Bloc until next upcall is ready

Rust System calls: command & allow

```
pub unsafe fn command(major: u32, minor: u32,  
    arg: isize) -> isize;
```

```
pub unsafe fn allow(major: u32, minor: u32,  
    slice: &[u8]) -> isize;
```

Rust System calls: subscribe

```
type ExternFn =  
    extern fn (usize, usize, usize, *const usize);  
  
pub unsafe fn subscribe(major: u32, minor: u32,  
    cb: ExternFn, ud: *const usize) -> isize {
```

Rust System calls: yieldk & yieldk_for


```
pub fn yieldk();
```

```
pub fn yieldk_for<F: Fn() -> bool>(cond: F) {  
    while !cond() { yieldk(); }  
}
```

Example: printing to the debug console

```
pub fn write(string: Box<[u8]>) {  
    let done: Cell<bool> = Cell::new(false);  
    syscalls::allow(DRIVER_NUM, 1, string);  
    syscalls::subscribe(DRIVER_NUM, 1, callback,  
        &done as *const _ as usize);  
    syscalls::command(DRIVER_NUM, 1, string.len());  
    yieldk_for(|| done.get())  
}  
  
extern fn callback(_: usize, _: usize, _: usize,  
    ud: *const usize) {  
    let done: &Cell<bool> = unsafe {  
        mem::transmute(ud)  
    };  
    done.set(true);  
}
```

Inter Process Communication (IPC)



Current Rust userland

Supported drivers

- ▶ Debug console
- ▶ Timer
- ▶ Sensors:
 - ▶ Accelerometer & magnetometer
 - ▶ Ambient light
 - ▶ Temperature
- ▶ IPC

Caveats

- ▶ No global variables allowed!
- ▶ Fixed code offset
- ▶ Blocking library calls only

Check your understanding

1. How does a process perform a blocking operation? Can you draw the flow of operations when a process calls `delay_ms(1000)`?
2. What is a Grant? How do processes interact with grants?
Hint: Think about memory exhaustion.

Hands-on: Write a BLE environment sensing application

1. Get a Rust application running on Hail
2. Periodically sample on-board sensors
3. Extend your app to report through the `ble-env-sense` service