

Tock Embedded OS Tutorial

SenSys 2018

Tock

A secure operating system for microcontrollers

- ▶ Kernel components in Rust
- ▶ Type-safe API for safe driver development
- ▶ Hardware isolated processes for application code

Microcontrollers

System-on-a-chip with integrated flash, SRAM, CPU and a bunch of hardware controllers.

Typically:

- ▶ Communication: UART, SPI, I2C, USB, CAN...
- ▶ External I/O: GPIO, external interrupt, ADC, DAC
- ▶ Timers: RTC, countdown timers

Maybe...

- ▶ Radio (Bluetooth, 15.4)
- ▶ Cryptographic accelerators
- ▶ Other specialized hardware...


Low Resource

- ▶ 10's of μA average power draw
- ▶ 10's of kB of RAM
- ▶ Moderate clock speeds

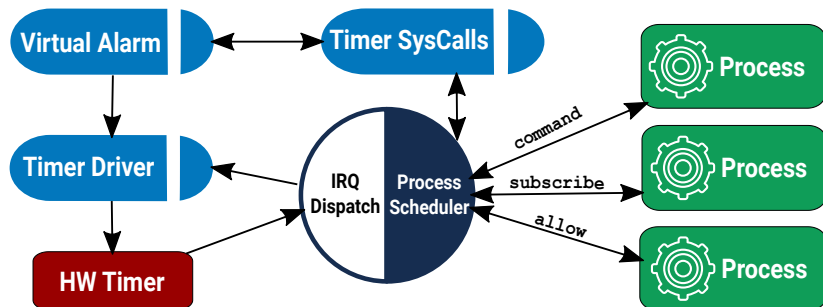
Use cases

- ▶ Security applications (e.g. authentication keys)
- ▶ Sensor networks
- ▶ Programmable wearables
- ▶ PC/phone peripherals
- ▶ Home/industrial automation
- ▶ Flight control

Two types of components: capsules and processes



Two types of scheduling: cooperative and preemptive




Agenda Today

09:30-10:40	Intro to Tock, Development Environment & Hardware
10:40-11:00	Coffee break
11:00-12:00	Hardware setup and installing apps
12:00-13:30	Lunch
13:30-15:20	Find and fix a real world bug
15:20-15:40	Coffee break
15:40-17:30	Choose your own adventure

Part 1: Hardware, tools, and development environment

Designed for TockOS
www.tockos.org

IMIX
Shane Leonard
v2.1 9/17
www.github.com/helena-project/imix



- ▶ Atmel SAM4L, Cortex-M4, 64 kB RAM, 256 kB flash
- ▶ Nordic NRF51 Bluetooth SoC
- ▶ 802.15.4 radio (6lowpan)
- ▶ Temperature, humidity, and light sensors
- ▶ 2 USBs (target USB + FTDI serial USB)
- ▶ 2 LEDs, 1 “user” button

Binaries on-board in flash

- ▶ 0x00000: **Bootloader**: Interact with Tockloader; load code
- ▶ 0x10000: **Kernel**
- ▶ 0x40000: **Processes**: Packed back-to-back

Tools

- ▶ `make`
- ▶ Rust/Cargo (Rust code → Cortex-M)
- ▶ `arm-none-eabi` (C → Cortex-M)
- ▶ `tockloader` to interact with imix and the bootloader

Tools: tockloader

Write a binary to a particular address in flash

```
$ tockloader flash --address 0x10000 \  
    target/thumbv7em-none-eabi/release/imix.bin
```

Program a process in Tock Binary Format¹:

```
$ tockloader install myapp.tab
```

Restart the board and connect to the debug console:

```
$ tockloader listen
```

¹TBFs are relocatable process binaries prefixed with headers like the package name.
.tab is a tarball of TBFs for different architectures as well as a metadata file for tockloader.

Check your understanding

Turn to the person next to you:

1. What kinds of binaries exist on a Tock board?

Hint: There are three, and only two can be programmed using `tockloader`.

2. What steps would you follow to program a process onto imix? What about to replace the kernel?

Answers

1. The three binaries are the serial bootloader, the kernel, and a series of processes. The bootloader can be used to load the kernel and processes, but cannot replace itself.
2. Use tockloader:
 - ▶ `tockloader install app.tab`
 - ▶ `tockloader flash --address 0x10000 imix-kernel.bin`

Hands-on: Set-up development environment

3. Compile and program the kernel
4. (Optional) Familiarize yourself with `tockloader` commands

- * ``uninstall``

- * ``list``

- * ``erase-apps``

5. (Optional) Add some other apps from the repo, like `blink` and `sensors`

▶ Head to <http://j2x.us/tock1> to get started!

▶ (<https://github.com/tock/tock/blob/tutorial-sensys-2018/doc/courses/sensys/environment.md>)

Part 2: User space programming

System calls

Tock supports five syscalls that applications use to interact with the kernel.

Call	Target	Description
command	Capsule	Invoke an operation on a capsule
allow	Capsule	Share memory with a capsule
subscribe	Capsule	Register an upcall
memop	Core	Modify memory break
yield	Core	Block until next upcall is ready

C System Calls: command & allow

// Start an operation

```
int command(u32 driver, u32 command, int arg1, int arg2);
```

// Share memory with the kernel

```
int allow(u32 driver, u32 allow, void* ptr, size_t size);
```

C System Calls: subscribe

```
// Callback function type  
typedef void (sub_cb)(int, int, int, void* userdata);  
  
// Register a callback with the kernel  
int subscribe(u32 driver,  
              u32 subscribe,  
              sub_cb cb,  
              void* userdata);
```

C System Calls: yield & yield_for

```
// Block until next callback
```

```
void yield(void);
```

```
// Block until a specific callback
```

```
void yield_for(bool *cond) {
```

```
    while (!*cond) {
```

```
        yield();
```

```
    }
```

```
}
```

Example: printing to the debug console

```
#define DRIVER_NUM_CONSOLE 0x0001

bool done = false;

static void putstr_cb(int x, int y, int z, void* ud) {
    done = true;
}

int putnstr(const char *str, size_t len) {
    allow(DRIVER_NUM_CONSOLE, 1, str, len);
    subscribe(DRIVER_NUM_CONSOLE, 1, putstr_cb, NULL);
    command(DRIVER_NUM_CONSOLE, 1, len, 0);
    yield_for(&done);

    return SUCCESS;
}
```

Hands-on: Write a simple application

3. Get an application running on imix
4. Print “Hello World” every second
5. Extend your app to sample on-board sensors

▶ Head to <http://j2x.us/tock2> to get started!

▶ (<https://github.com/tock/tock/blob/tutorial-sensys-2018/doc/courses/sensys/application.md>)

Part 3: The kernel

Trusted Computing Base (unsafe allowed)

- ▶ Hardware Abstraction Layer
- ▶ Board configuration
- ▶ Event & Process scheduler
- ▶ Rust core library
- ▶ Core Tock primitives

arch/

chips/

kernel/

Capsules (unsafe not allowed)

- ▶ Virtualization
- ▶ Peripheral drivers
- ▶ Communication protocols (IP, USB, etc)
- ▶ Application logic

capsules/

Constraints

Small isolation units

Breaking a monolithic component into smaller ones should have low/no cost

Avoid memory exhaustion in the kernel

No heap. Everything is allocated statically.

Low communication overhead

Communicating between components as cheap as an internal function call.
Ideally inlined.

Event-driven execution model

```
pub fn main<P, C>(platform: &P, chip: &mut C,
                  processes: &mut [Process]) {
    loop {
        chip.service_pending_interrupts();
        for (i, p) in processes.iter_mut().enumerate() {
            sched::do_process(platform, chip, process);
        }


        if !chip.has_pending_interrupts() {
            chip.prepare_for_sleep();
            support::wfi();
        }
    }
}
```

Event-driven execution model

```
fn service_pending_interrupts(&mut self) {  
    while let Some(interrupt) = get_interrupt() {  
        match interrupt {  
            ASTALARM => ast::AST.handle_interrupt(),  
            USART0 => usart::USART0.handle_interrupt(),  
            USART1 => usart::USART1.handle_interrupt(),  
            USART2 => usart::USART2.handle_interrupt(),  
            ...  
        }  
    }  
}
```

Event-driven execution model

```
impl Ast {  
    pub fn handle_interrupt(&self) {  
        self.clear_alarm();  
        self.callback.get().map(|cb| { cb.fired(); });  
    }  
}  
  
impl time::Client for MuxAlarm {  
    fn fired(&self) {  
        for cur in self.virtual_alarms.iter() {  
            if cur.should_fire() {  
                cur.armed.set(false);  
                self.enabled.set(self.enabled.get() - 1);  
                cur.fired();  
            }  
        }  
    }  
}
```



Check your understanding

Turn to the person next to you:

1. What are Tock kernel components called?
2. Is the kernel scheduled cooperatively or preemptively? What happens if a capsule performs a very long computation?
3. How is a hardware interrupt handled in the kernel?

Answers

1. Tock kernel components are called “capsules”
2. The kernel is scheduled cooperatively by capsules calling methods on each other. If a capsule performs a very long computation it might prevent other capsules from running or cause them to miss events.
3. Hardware interrupts are scheduled to run when capsules next yield. If a process is running when a hardware event happens, the hardware event will be immediately handled.

Hands-on: Write and add a capsule to the kernel

4. Read the imix boot sequence in `boards/imix/src/main.rs`
5. Write a new capsule that prints “Hello World” to the debug console.
6. Extend your capsule to print “Hello World” every second
7. Extend your capsule to print light readings every second
8. Extra credit

▶ Head to <http://j2x.us/tock3> to get started!

▶ (<https://github.com/tock/tock/blob/tutorial-sensys-2018/doc/courses/sensys/capsule.md>)

Part 3: Deliver for the Client

Debugging in a Multi-app Setting

- ▶ Multiprogramming -> multiple things can go wrong
- ▶ Multiprogramming *enables* better debugging facilities
 - ▶ Monitor individual application state
 - ▶ Disable only faulty applications
 - ▶ Replace only parts of the system

The Process Console on imix

Keeps track of system calls and timeslice expirations for each process.

Provides basic debugging facilities over UART:

- ▶ `status`
- ▶ `list`
- ▶ `stop [app_name]`
- ▶ `start [app_name]`

Our task

- ▶ Fix a “deployed” with two mysteriously named processes:
 - ▶ app1 & app2
- ▶ Networked over UDP/6LoWPAN
 - ▶ One sends button presses
 - ▶ Another sends periodic temperature, humidity and light data
- ▶ Functional, but seem to be draining battery!

Find and fix the problem!

Stay in touch!

<https://www.tockos.org>

<https://github.com/tock/tock>

tock-dev@googlegroups.com
